

Lab #6: eBPF

Kotaiba Alachkar 8-)

Preparation

```
sudo apt install vagrant virtualbox wget
sudo modprobe vboxdrv vboxnetadp vboxnetflt vboxpci
wget http://amiens.studlab.os3.nl/an2018/lab-ebpf/an-ebpf.tgz
tar xf an-ebpf.tgz
cd an-ebpf
vagrant up
vagrant ssh
sudo su
cd /vagrant
cd src
make
```

The output command should state **'SUCCESS'** in the last line:

```
+ tc qdisc del dev eth1 clsact
+ set +x
+ tc qdisc add dev eth1 clsact
+ tc filter add dev eth1 ingress prio 1 handle 1 bpf da obj tcbpf1_kern.o
sec classifier
+ set +x
SUCCESS
```

Task 1: Compiling and usage of sample eBPF program (2 points)

```
/*
 * AN2018 lab6 ebpf
 *
 * This is a modified verison of linux/samples/bpf/tcbpf1_kern.c
 *
 * it includes bpf_debug.h to output debug information that can be
 * read using tools/bpf-trace.
 *
 * ifindex can be obtained using: ip link show | cut -c 1
 *
 */

#define KBUILD_MODNAME "foo"
#include <uapi/linux/bpf.h>
```

```

#include <uapi/linux/if_ether.h>
#include <uapi/linux/if_packet.h>
#include <uapi/linux/ip.h>
#include <uapi/linux/in.h>
#include <uapi/linux/tcp.h>
#include <uapi/linux/udp.h>
#include <uapi/linux/filter.h>
#include <uapi/linux/pkt_cls.h>
#include "bpf_helpers.h"
#include "bpf_debug.h"

/* compiler workaround */
#define _htonl __builtin_bswap32

static inline void set_dst_mac(struct __sk_buff *skb, char *mac)
{
    //sets destination mac
    bpf_skb_store_bytes(skb, 0, mac, ETH_ALEN, 1);
}

#define IP_CSUM_OFF (ETH_HLEN + offsetof(struct iphdr, check))
#define TOS_OFF (ETH_HLEN + offsetof(struct iphdr, tos))

static inline void set_ip_tos(struct __sk_buff *skb, __u8 new_tos)
{
    //sets tos, and recalculate checksum
    __u8 old_tos = load_byte(skb, TOS_OFF);

    bpf_l3_csum_replace(skb, IP_CSUM_OFF, htons(old_tos), htons(new_tos),
2);
    bpf_skb_store_bytes(skb, TOS_OFF, &new_tos, sizeof(new_tos), 0);
}

#define IP_SRC_OFF (ETH_HLEN + offsetof(struct iphdr, saddr))

#define TCP_CSUM_OFF (ETH_HLEN + sizeof(struct iphdr) + offsetof(struct
tcphdr, check))
#define UDP_CSUM_OFF (ETH_HLEN + sizeof(struct iphdr) + offsetof(struct
udphdr, check))

#define IS_PSEUDO 0x10

static inline void set_tcp_ip_src(struct __sk_buff *skb, __u32 new_ip) {
    //set source address for tcp and recalculates checksum
    __u32 old_ip = _htonl(load_word(skb, IP_SRC_OFF));

    bpf_l4_csum_replace(skb, TCP_CSUM_OFF, old_ip, new_ip, IS_PSEUDO |
sizeof(new_ip));
    bpf_l3_csum_replace(skb, IP_CSUM_OFF, old_ip, new_ip, sizeof(new_ip));
    bpf_skb_store_bytes(skb, IP_SRC_OFF, &new_ip, sizeof(new_ip), 0);
}

```

```

}

#define TCP_DPORT_OFF (ETH_HLEN + sizeof(struct iphdr) + offsetof(struct
tcphdr, dest))
static inline void set_tcp_dest_port(struct __sk_buff *skb, __u16 new_port)
{
    //set destination port for tcp and recalculates checksum
    __u16 old_port = htons(load_half(skb, TCP_DPORT_OFF));

    bpf_l4_csum_replace(skb, TCP_CSUM_OFF, old_port, new_port,
sizeof(new_port));
    bpf_skb_store_bytes(skb, TCP_DPORT_OFF, &new_port, sizeof(new_port), 0);
}

#define UDP_DPORT_OFF (ETH_HLEN + sizeof(struct iphdr) + offsetof(struct
udphdr, dest))
static inline void set_udp_dest_port(struct __sk_buff *skb, __u16 new_port)
{
    //set destinationport for udp and recalculates checksum
    __u16 old_port = htons(load_half(skb, UDP_DPORT_OFF));

    bpf_l4_csum_replace(skb, UDP_CSUM_OFF, old_port, new_port,
sizeof(new_port));
    bpf_skb_store_bytes(skb, UDP_DPORT_OFF, &new_port, sizeof(new_port), 0);
}

static inline void set_udp_ip_src(struct __sk_buff *skb, __u32 new_ip) {
    //set source address for udp and recalculats checksum
    __u32 old_ip = _htonl(load_word(skb, IP_SRC_OFF));

    bpf_l4_csum_replace(skb, UDP_CSUM_OFF, old_ip, new_ip, IS_PSEUDO |
sizeof(new_ip));
    bpf_l3_csum_replace(skb, IP_CSUM_OFF, old_ip, new_ip, sizeof(new_ip));
    bpf_skb_store_bytes(skb, IP_SRC_OFF, &new_ip, sizeof(new_ip), 0);
}

static inline __u32 ip(__u8 o1, __u8 o2, __u8 o3, __u8 o4){
    /* takes the four octets of a ip address and calculates the hex version
    * in network byte order*/
    __u32 result = 0;
    result += o4 * 256*256*256;
    result += o3 * 256*256;
    result += o2 * 256;
    result += o1;
    return result;
}

static inline void print_ip( __u32 ip){
    /* prints 'human readable' ipv4 address using bpf_debug
    * the function writes two output lines since kprint is limited
    * to three arguments */

```

```

__u8 o1 = (__u8) ip;
__u8 o2 = (__u8) (ip >> 8);
__u8 o3 = (__u8) (ip >> 16);
__u8 o4 = (__u8) (ip >> 24);
bpf_debug("ip(1) %d.%d.x.x\n", o1, o2);
bpf_debug("ip(2) x.x.%d.%d\n", o3, o4);
}

SEC("rewrite_tcp")
int _rewrite_tcp(struct __sk_buff *skb)
{
    __u8 proto = load_byte(skb, ETH_HLEN + offsetof(struct iphdr,
protocol));
    __u8 ifindex = 255;

    if (proto == IPPROTO_TCP) {
        __u8 old_tos = load_byte(skb, TOS_OFF);
        __u16 old_port = load_half(skb, TCP_DPORT_OFF);
        __u32 old_ip = _htonl(load_word(skb, IP_SRC_OFF));

        //set_udp_ip_src(skb, 0xA010101); //1.1.1.10
        //set_udp_ip_src(skb, 0xFE02A8C0); //192.168.2.254
        //set_udp_ip_src(skb, 0xA0A0A0A0); //10.10.10.10

        set_udp_ip_src(skb, ip(10,10,10,10)); //10.10.10.10
        set_ip_tos(skb, 8);
        set_tcp_dest_port(skb, htons(8000));

        __u8 new_tos = load_byte(skb, TOS_OFF);
        __u16 new_port = load_half(skb, TCP_DPORT_OFF);
        __u32 new_ip = _htonl(load_word(skb, IP_SRC_OFF));

        bpf_debug("rewrote tos %d -> %d\n", old_tos, new_tos);
        bpf_debug("rewrote src_ip %x -> %x\n", old_ip, new_ip);
        bpf_debug("<old_dest_ip>\n");
        print_ip(old_ip);
        bpf_debug("</old_dest_ip>\n");
        bpf_debug("<new_dest_ip>\n");
        print_ip(new_ip);
        bpf_debug("</new_dest_ip>\n");
        bpf_debug("rewrote dst_port %d -> %d\n", old_port, new_port);

        if (ifindex != 255) {
            return bpf_clone_redirect(skb, ifindex, 0);
        }
    }
    return BPF_OK;
}

SEC("rewrite_udp")
int _rewrite_udp(struct __sk_buff *skb) {

```

```

__u8 proto = load_byte(skb, ETH_HLEN + offsetof(struct iphdr,
protocol));
__u8 ifindex = 255;

if (proto == IPPROTO_UDP) {
    __u8 old_tos = load_byte(skb, TOS_OFF);
    __u16 old_port = load_half(skb, UDP_DPORT_OFF);
    __u32 old_ip = _htonl(load_word(skb, IP_SRC_OFF));

    // rewrite packet
    //set_udp_ip_src(skb, 0xA010101); //1.1.1.10
    //set_udp_ip_src(skb, 0xFE02A8C0); //192.168.2.254
    //set_udp_ip_src(skb, 0xA0A0A0A0); //10.10.10.10

    set_udp_ip_src(skb, ip(192,168,2,254)); //192.168.2.254
    set_ip_tos(skb, 8);
    set_udp_dest_port(skb, htons(8000));

    __u8 new_tos = load_byte(skb, TOS_OFF);
    __u16 new_port = load_half(skb, UDP_DPORT_OFF);
    __u32 new_ip = _htonl(load_word(skb, IP_SRC_OFF));

    bpf_debug("rewrote tos %d -> %d\n", old_tos, new_tos);
    bpf_debug("rewrote src_ip %x -> %x\n", old_ip, new_ip);
    bpf_debug("<old_dest_ip>\n");
    print_ip(old_ip);
    bpf_debug("</old_dest_ip>\n");
    bpf_debug("<new_dest_ip>\n");
    print_ip(new_ip);
    bpf_debug("</new_dest_ip>\n");
    bpf_debug("rewrote dst_port %d -> %d\n", old_port, new_port);

    if (ifindex != 255) {
        return bpf_clone_redirect(skb, ifindex, 0);
    }
}
return BPF_OK;
}

SEC("task4")
int _task4(struct __sk_buff *skb) {
    /* Modify this function to filter certain packets.
     * Use the code above and linux-4.15/include/uapi/linux/bpf.h
     * as a reference. */

    return BPF_OK;
}

char _license[] SEC("license") = "GPL";

```

Q1.1 Look at the code section SEC("rewrite_tcp") and explain the code within. Be specific, explain the role of used function parameters and the purpose of returned values.

SEC("rewrite_tc") code snippet:

```
SEC("rewrite_tcp")
int _rewrite_tcp(struct __sk_buff *skb)
{
    __u8 proto = load_byte(skb, ETH_HLEN + offsetof(struct iphdr,
protocol));
    __u8 ifindex = 255;

    if (proto == IPPROTO_TCP) {
        __u8 old_tos = load_byte(skb, TOS_OFF);
        __u16 old_port = load_half(skb, TCP_DPORT_OFF);
        __u32 old_ip = _htonl(load_word(skb, IP_SRC_OFF));

        //set_udp_ip_src(skb, 0xA010101); //1.1.1.10
        //set_udp_ip_src(skb, 0xFE02A8C0); //192.168.2.254
        //set_udp_ip_src(skb, 0xA0A0A0A0); //10.10.10.10

        set_udp_ip_src(skb, ip(10,10,10,10)); //10.10.10.10
        set_ip_tos(skb, 8);
        set_tcp_dest_port(skb, htons(8000));

        __u8 new_tos = load_byte(skb, TOS_OFF);
        __u16 new_port = load_half(skb, TCP_DPORT_OFF);
        __u32 new_ip = _htonl(load_word(skb, IP_SRC_OFF));

        bpf_debug("rewrote tos %d -> %d\n", old_tos, new_tos);
        bpf_debug("rewrote src_ip %x -> %x\n", old_ip, new_ip);
        bpf_debug("<old_dest_ip>\n");
        print_ip(old_ip);
        bpf_debug("</old_dest_ip>\n");
        bpf_debug("<new_dest_ip>\n");
        print_ip(new_ip);
        bpf_debug("</new_dest_ip>\n");
        bpf_debug("rewrote dst_port %d -> %d\n", old_port, new_port);

        if (ifindex != 255) {
            return bpf_clone_redirect(skb, ifindex, 0);
        }
    }
    return BPF_OK;
}
```

The SEC("rewrite_tcp") is changing the SourceIP, DstPort and TOS parameters within the packet using a predefined function if the protocol value equals to the one given in the IPPROTO_TCP constant value. In addition to that if the ifindex is not the default value "255", the code will redirect the packet

to a new interface, specified as ifindex.

The input of the function is buffer skb. load_byte loads one byte on a specific offset from buffer skb. load_half loads two bytes on a specific offset from buffer skb. load_word loads four bytes on a specific offset from buffer skb.

The function returns an integer. It contains multiple definitions of unsigned integers:

- u8: Unsigned 8-bit integer * u16: Unsigned 16-bit integer
- u32: Unsigned 32-bit integer ===== Q2.2 - Compile the program using tools/bpf-compile wrapper: tools/bpf-compile <filename>. Apply the program to eth1 interface using: ./tools/bpf-tc eth1 bpf_sample.o rewrite_tcp Verify that the object file is loaded and provide the output of the above command. Hint: use tc. ===== Compile the program using tools/bpf-compile wrapper:

```
root@archlinux:/vagrant/src# ./tools/bpf-compile bpf_sample.c ++ sed -e 's/\.c$/\.o/' ++ basename bpf_sample.c ++ echo bpf_sample.c + OBJ_FILENAME=bpf_sample.o + LLC=llc + CLANG=clang + ARCH=x86 + ARCH_FULL=x86_64-linux-gnu ++ uname -r + HEADERS=/usr/src/linux-headers-4.15.3-1-ARCH + KERNEL_MASTER=/vagrant/linux-4.15 + llc -march=bpf -filetype=obj -o bpf_sample.o + exec clang -nostdinc -isystem /usr/lib/gcc/x86_64-linux-gnu/6/include -I/vagrant/linux-4.15/arch/x86/include -I/vagrant/linux-4.15/arch/x86/include/generated -I/vagrant/linux-4.15/include -I/usr/lib/clang/5.0.1/include -I/vagrant/linux-4.15/arch/x86/include/uapi -I/vagrant/linux-4.15/arch/x86/include/generated/uapi -I/vagrant/linux-4.15/include/uapi -I/vagrant/linux-4.15/include/generated/uapi -include/vagrant/linux-4.15/include/linux/kconfig.h -I/vagrant/linux-4.15/samples/bpf -I/vagrant/linux-4.15/tools/testing/selftests/bpf/ -I/vagrant/include -DKERNEL -DASM_SYSREG_H -Wno-unused-value -Wno-pointer-sign -fno-stack-protector -Wno-compare-distinct-pointer-types -Wno-gnu-variable-sized-type-not-at-end -Wno-address-of-packed-member -Wno-tautological-compare -Wno-unknown-warning-option -O2 -emit-llvm -c bpf_sample.c -o -
```

</code>

Apply the program to eth1 interface:

```
root@archlinux:/vagrant/src# ./tools/bpf-tc eth1 bpf_sample.o rewrite_tcp
+ tc qdisc del dev eth1 clsact
+ set +x
+ tc qdisc add dev eth1 clsact
+ tc filter add dev eth1 ingress prio 1 handle 1 bpf da obj bpf_sample.o sec
rewrite_tcp
+ set +x
```

Verify the object file:

```
root@archlinux:/vagrant/src# tc filter show dev eth1 ingress
filter protocol all pref 1 bpf chain 0
filter protocol all pref 1 bpf chain 0 handle 0x1 bpf_sample.o:[rewrite_tcp]
direct-action not_in_hw id 15 tag e7983f733cf41b3f jited
```

Q3.3 - Illustrate the functionality realized by the attached program. Use tools

such as ping or nc to generate sample packets and tcpdump for your packet traces.

Create a dummy net interface:

```
root@archlinux:/vagrant/src# ip link add dummy0 type dummy
root@archlinux:/vagrant/src# ip link set dev dummy0 up

root@archlinux:/vagrant/src# ip link show dummy0
4: dummy0: <BROADCAST,NOARP,UP,LOWER_UP> mtu 1500 qdisc noqueue state
UNKNOWN mode DEFAULT group default qlen 1000
    link/ether 72:7b:b7:82:94:4f brd ff:ff:ff:ff:ff:ff
```

I changed the interface index in bpf_sample.c to 4 since dummy0 is in index 4:

```
__u8 ifindex = 4;
```

Now, recompile again:

```
root@archlinux:/vagrant/src# ../tools/bpf-compile bpf_sample.c
++ sed -e 's/\.c$/\.o/'
+++ basename bpf_sample.c
++ echo bpf_sample.c
+ OBJ_FILE=bpf_sample.o
+ LLC=llc
+ CLANG=clang
+ ARCH=x86
+ ARCH_FULL=x86_64-linux-gnu
++ uname -r
+ HEADERS=/usr/src/linux-headers-4.15.3-1-ARCH
+ KERNEL_MASTER=/vagrant/linux-4.15
+ exec clang -nostdinc -isystem /usr/lib/gcc/x86_64-linux-gnu/6/include -
I/vagrant/linux-4.15/arch/x86/include -
I/vagrant/linux-4.15/arch/x86/include/generated -
I/vagrant/linux-4.15/include -I/usr/lib/clang/5.0.1/include -
I/vagrant/linux-4.15/arch/x86/include/uapi -
I/vagrant/linux-4.15/arch/x86/include/generated/uapi -
I/vagrant/linux-4.15/include/uapi -
I/vagrant/linux-4.15/include/generated/uapi -
include/vagrant/linux-4.15/include/linux/kconfig.h -
I/vagrant/linux-4.15/samples/bpf -
I/vagrant/linux-4.15/tools/testing/selftests/bpf/ -I/vagrant/include -
D__KERNEL__ -D__ASM_SYSREG_H -Wno-unused-value -Wno-pointer-sign -fno-stack-
protector -Wno-compare-distinct-pointer-types -Wno-gnu-variable-sized-type-
not-at-end -Wno-address-of-packed-member -Wno-tautological-compare -Wno-
unknown-warning-option -O2 -emit-llvm -c bpf_sample.c -o -
+ llc -march=bpf -filetype=obj -o bpf_sample.o

root@archlinux:/vagrant/src# ../tools/bpf-tc eth1 bpf_sample.o rewrite_tcp
+ tc qdisc del dev eth1 clsact
```



```
+ set +x
+ tc qdisc add dev eth1 clsact
+ tc filter add dev eth1 ingress prio 1 handle 1 bpf da obj bpf_sample.o sec
rewrite_tcp
+ set +x
```

Setup a tcpdump on dummy0 and use nc to the vagrant box using "nc 192.168.2.100 80":

On my server:

```
root@bristol:~# nc 192.168.2.100 80
```

On vagrant:

```
tcpdump: listening on dummy0, link-type EN10MB (Ethernet), capture size
262144 bytes
07:44:06.145936 IP (tos 0x8, ttl 64, id 40173, offset 0, flags [DF], proto
TCP (6), length 60)
    10.10.10.10.37748 > 192.168.2.100.8000: Flags [S], cksum 0x3ef9
(correct), seq 734858136, win 29200, options [mss 1460,sackOK,TS val
332251631 ecr 0,nop,wscale 7], length 0
```

The ToS value was changed to 8 and the source IP-address to "10.10.10.10".

Task 2: Writing eBPF program: traffic firewalling (3 points)

Q2.1 Implement your filtering program in the sample by creating your own, it can be as simple as accepting only specific protocols (e.g. TCP+IPv4 only) or verifying port numbers and IP addresses. Use the section SEC("task4"). Use tcpdump to verify that it works.

Only capture a non-80 traffic:

```
SEC("task4")
int _task4(struct __sk_buff *skb) {
    /* Modify this function to filter certain packets.
     * Use the code above and linux-4.15/include/uapi/linux/bpf.h
     * as a reference. */

    if (oldPort == 80) {
        return BPF_DROP;
    } else {
        return bpf_clone_redirect(skb, ifindex, 0);
    }
}
```

Task 3: Filtering performance measurements (3 points)

Compiling:

```
root@archlinux:/vagrant# tools/bpf-compile src/ip_filter_w_map.c
++ sed -e 's/\.c$/\.o/'
+++ basename src/ip_filter_w_map.c
++ echo ip_filter_w_map.c
+ OBJ_FILENAME=ip_filter_w_map.o
+ LLC=llc
+ CLANG=clang
+ ARCH=x86
+ ARCH_FULL=x86_64-linux-gnu
++ uname -r
+ HEADERS=/usr/src/linux-headers-4.15.3-1-ARCH
+ KERNEL_MASTER=/vagrant/linux-4.15
+ exec clang -nostdinc -isystem /usr/lib/gcc/x86_64-linux-gnu/6/include -
I/vagrant/linux-4.15/arch/x86/include -
I/vagrant/linux-4.15/arch/x86/include/generated -
I/vagrant/linux-4.15/include -I/usr/lib/clang/5.0.1/include -
I/vagrant/linux-4.15/arch/x86/include/uapi -
I/vagrant/linux-4.15/arch/x86/include/generated/uapi -
I/vagrant/linux-4.15/include/uapi -
I/vagrant/linux-4.15/include/generated/uapi -
include/vagrant/linux-4.15/include/linux/kconfig.h -
I/vagrant/linux-4.15/samples/bpf -
I/vagrant/linux-4.15/tools/testing/selftests/bpf/ -I/vagrant/include -
D__KERNEL__ -D__ASM_SYSREG_H -Wno-unused-value -Wno-pointer-sign -fno-stack-
protector -Wno-compare-distinct-pointer-types -Wno-gnu-variable-sized-type-
not-at-end -Wno-address-of-packed-member -Wno-tautological-compare -Wno-
unknown-warning-option -O2 -emit-llvm -c src/ip_filter_w_map.c -o -
+ llc -march=bpf -filetype=obj -o ip_filter_w_map.o

root@archlinux:/vagrant# tools/bpf-tc eth1 ip_filter_w_map.o classifier
+ tc qdisc del dev eth1 clsact
+ set +x
+ tc qdisc add dev eth1 clsact
+ tc filter add dev eth1 ingress prio 1 handle 1 bpf da obj
ip_filter_w_map.o sec classifier
Note: 8 bytes struct bpf_elf_map fixup performed due to size mismatch!
+ set +x
```

Verification of bpf-map:

```
root@archlinux:/vagrant# bpf-map info /sys/fs/bpf/tc/globals/ddos
Type:      Hash
Key size:  4
Value size: 8
Max entries: 11000
```

Flags: 0x0

Remove eBPF object file and check again:

```
root@archlinux:/vagrant# tc filter show ingress dev eth1
filter protocol all pref 1 bpf chain 0
filter protocol all pref 1 bpf chain 0 handle 0x1
ip_filter_w_map.o:[classifier] direct-action not_in_hw id 17 tag
33fbcfbd5e591bb2 jited
root@archlinux:/vagrant# tc qdisc del dev eth1 clsact
root@archlinux:/vagrant# tc filter show ingress dev eth1
```

Q3.1 - Execute iptables -t raw -i eth1 -A PREROUTING -j NOTRACK on the vagrant VM to prevent 'conntrack table' space exhaustion. Start an iperf3 server on the VM and run iperf3 -c 192.168.2.100 -t 70 -O 10 on the host machine to test the bandwidth between the host and VM. First, run it as is to see the raw performance. Also on the host system, run hping3 --rand-source 192.168.2.100 --faster command to start generating DDoS traffic; leave hping3 running in the background. If hping3 kills your connection, change --faster to -i u1000 and change the number to increase or decrease the sending rate of hping3. The goal is to see slight degradation of iperf speeds (caused by hping3) without any rules applied. Show what you did, and that a hping3 is properly tuned.

First, Execute:

```
root@archlinux:/vagrant# iptables -t raw -i eth1 -A PREROUTING -j NOTRACK
root@archlinux:/vagrant# iperf3 -s
-----
Server listening on TCP port 5001
TCP window size: 85.3 KByte (default)
-----
```

On my server:

```
root@bristol:/home/kotaiba# iperf3 -c 192.168.2.100 -t 70 -O 10
Connecting to host 192.168.2.100, port 5201
[ 4] local 192.168.2.1 port 49758 connected to 192.168.2.100 port 5201
[ ID] Interval          Transfer      Bandwidth      Retr  Cwnd
[ 4]  0.00-1.00    sec    228 MBytes   1.91 Gbits/sec   96   249 KBytes
(omitted)
[ 4]  1.00-2.00    sec    231 MBytes   1.93 Gbits/sec  161   387 KBytes
(omitted)
[ 4]  2.00-3.00    sec    227 MBytes   1.91 Gbits/sec   24   437 KBytes
(omitted)
[ 4]  3.00-4.00    sec    230 MBytes   1.93 Gbits/sec   25   450 KBytes
(omitted)
[ 4]  4.00-5.00    sec    240 MBytes   2.01 Gbits/sec   62   317 KBytes
```

```
(omitted)
[ 4] 5.00-6.00 sec 247 MBytes 2.07 Gbits/sec 43 411 KBytes
(omitted)
[ 4] 6.00-7.00 sec 249 MBytes 2.09 Gbits/sec 53 273 KBytes
(omitted)
[ 4] 7.00-8.00 sec 245 MBytes 2.06 Gbits/sec 94 379 KBytes
(omitted)
[ 4] 8.00-9.00 sec 247 MBytes 2.07 Gbits/sec 125 363 KBytes
(omitted)
[ 4] 9.00-10.00 sec 219 MBytes 1.84 Gbits/sec 60 424 KBytes
(omitted)
[ 4] 0.00-1.00 sec 241 MBytes 2.02 Gbits/sec 49 468 KBytes
[ 4] 1.00-2.00 sec 221 MBytes 1.86 Gbits/sec 24 452 KBytes
[ 4] 2.00-3.00 sec 225 MBytes 1.88 Gbits/sec 52 455 KBytes
[ 4] 3.00-4.00 sec 225 MBytes 1.89 Gbits/sec 12 464 KBytes
[ 4] 4.00-5.00 sec 227 MBytes 1.91 Gbits/sec 57 406 KBytes
[ 4] 5.00-6.00 sec 221 MBytes 1.85 Gbits/sec 11 396 KBytes
[ 4] 6.00-7.00 sec 229 MBytes 1.92 Gbits/sec 181 354 KBytes
[ 4] 7.00-8.00 sec 223 MBytes 1.87 Gbits/sec 108 404 KBytes
[ 4] 8.00-9.00 sec 219 MBytes 1.84 Gbits/sec 143 427 KBytes
[ 4] 9.00-10.00 sec 224 MBytes 1.88 Gbits/sec 108 368 KBytes
[ 4] 10.00-11.00 sec 228 MBytes 1.91 Gbits/sec 118 392 KBytes
[ 4] 11.00-12.00 sec 223 MBytes 1.87 Gbits/sec 12 417 KBytes
[ 4] 12.00-13.00 sec 224 MBytes 1.88 Gbits/sec 123 444 KBytes
[ 4] 13.00-14.00 sec 244 MBytes 2.05 Gbits/sec 29 433 KBytes
[ 4] 14.00-15.00 sec 246 MBytes 2.06 Gbits/sec 109 499 KBytes
[ 4] 15.00-16.00 sec 240 MBytes 2.02 Gbits/sec 90 427 KBytes
[ 4] 16.00-17.00 sec 227 MBytes 1.91 Gbits/sec 55 423 KBytes
[ 4] 17.00-18.00 sec 239 MBytes 2.00 Gbits/sec 107 378 KBytes
[ 4] 18.00-19.00 sec 226 MBytes 1.89 Gbits/sec 17 399 KBytes
[ 4] 19.00-20.00 sec 220 MBytes 1.85 Gbits/sec 54 329 KBytes
[ 4] 20.00-21.00 sec 226 MBytes 1.89 Gbits/sec 32 402 KBytes
[ 4] 21.00-22.00 sec 221 MBytes 1.85 Gbits/sec 75 406 KBytes
[ 4] 22.00-23.00 sec 228 MBytes 1.92 Gbits/sec 17 433 KBytes
[ 4] 23.00-24.00 sec 226 MBytes 1.89 Gbits/sec 149 362 KBytes
[ 4] 24.00-25.00 sec 233 MBytes 1.95 Gbits/sec 88 314 KBytes
[ 4] 25.00-26.00 sec 234 MBytes 1.96 Gbits/sec 168 352 KBytes
[ 4] 26.00-27.00 sec 232 MBytes 1.95 Gbits/sec 31 352 KBytes
[ 4] 27.00-28.00 sec 224 MBytes 1.88 Gbits/sec 15 467 KBytes
[ 4] 28.00-29.00 sec 224 MBytes 1.88 Gbits/sec 11 471 KBytes
[ 4] 29.00-30.00 sec 221 MBytes 1.85 Gbits/sec 54 477 KBytes
[ 4] 30.00-31.00 sec 229 MBytes 1.92 Gbits/sec 103 436 KBytes
[ 4] 31.00-32.00 sec 230 MBytes 1.93 Gbits/sec 66 396 KBytes
[ 4] 32.00-33.00 sec 235 MBytes 1.97 Gbits/sec 61 284 KBytes
[ 4] 33.00-34.00 sec 242 MBytes 2.03 Gbits/sec 63 313 KBytes
[ 4] 34.00-35.00 sec 246 MBytes 2.06 Gbits/sec 157 370 KBytes
[ 4] 35.00-36.00 sec 229 MBytes 1.92 Gbits/sec 61 444 KBytes
[ 4] 36.00-37.00 sec 235 MBytes 1.97 Gbits/sec 47 395 KBytes
[ 4] 37.00-38.00 sec 228 MBytes 1.91 Gbits/sec 185 304 KBytes
[ 4] 38.00-39.00 sec 225 MBytes 1.89 Gbits/sec 45 335 KBytes
[ 4] 39.00-40.00 sec 219 MBytes 1.84 Gbits/sec 59 332 KBytes
```

```

[ 4] 40.00-41.00 sec 229 MBytes 1.92 Gbits/sec 41 342 KBytes
[ 4] 41.00-42.00 sec 229 MBytes 1.92 Gbits/sec 19 354 KBytes
[ 4] 42.00-43.00 sec 228 MBytes 1.91 Gbits/sec 108 400 KBytes
[ 4] 43.00-44.00 sec 222 MBytes 1.86 Gbits/sec 61 441 KBytes
[ 4] 44.00-45.00 sec 228 MBytes 1.91 Gbits/sec 42 329 KBytes
[ 4] 45.00-46.00 sec 227 MBytes 1.91 Gbits/sec 24 304 KBytes
[ 4] 46.00-47.00 sec 224 MBytes 1.88 Gbits/sec 114 409 KBytes
[ 4] 47.00-48.00 sec 235 MBytes 1.97 Gbits/sec 197 389 KBytes
[ 4] 48.00-49.00 sec 236 MBytes 1.98 Gbits/sec 87 430 KBytes
[ 4] 49.00-50.00 sec 225 MBytes 1.89 Gbits/sec 60 423 KBytes
[ 4] 50.00-51.00 sec 222 MBytes 1.86 Gbits/sec 59 443 KBytes
[ 4] 51.00-52.00 sec 222 MBytes 1.86 Gbits/sec 93 368 KBytes
[ 4] 52.00-53.00 sec 236 MBytes 1.98 Gbits/sec 68 373 KBytes
[ 4] 53.00-54.00 sec 237 MBytes 1.99 Gbits/sec 83 392 KBytes
[ 4] 54.00-55.00 sec 229 MBytes 1.92 Gbits/sec 30 420 KBytes
[ 4] 55.00-56.00 sec 229 MBytes 1.92 Gbits/sec 166 383 KBytes
[ 4] 56.00-57.00 sec 219 MBytes 1.84 Gbits/sec 40 460 KBytes
[ 4] 57.00-58.00 sec 222 MBytes 1.86 Gbits/sec 45 458 KBytes
[ 4] 58.00-59.00 sec 230 MBytes 1.93 Gbits/sec 210 303 KBytes
[ 4] 59.00-60.00 sec 230 MBytes 1.93 Gbits/sec 57 378 KBytes
[ 4] 60.00-61.00 sec 219 MBytes 1.84 Gbits/sec 57 389 KBytes
[ 4] 61.00-62.00 sec 240 MBytes 2.01 Gbits/sec 44 291 KBytes
[ 4] 62.00-63.00 sec 223 MBytes 1.87 Gbits/sec 13 434 KBytes
[ 4] 63.00-64.00 sec 222 MBytes 1.86 Gbits/sec 29 457 KBytes
[ 4] 64.00-65.00 sec 218 MBytes 1.83 Gbits/sec 29 355 KBytes
[ 4] 65.00-66.00 sec 231 MBytes 1.94 Gbits/sec 57 455 KBytes
[ 4] 66.00-67.00 sec 228 MBytes 1.91 Gbits/sec 81 395 KBytes
^C[ 4] 67.00-67.48 sec 107 MBytes 1.89 Gbits/sec 9 385 KBytes
- - - - -
[ ID] Interval          Transfer      Bandwidth      Retr
[ 4] 0.00-67.48 sec 15.0 GBytes 1.92 Gbits/sec 4869
sender
[ 4] 0.00-67.48 sec 0.00 Bytes 0.00 bits/sec receiver
iperf3: interrupt - the client has terminated

```

With hping in the background, the following output are generated:

```

root@bristol:/home/kotaiba# iperf3 -c 192.168.2.100 -t 70 -0 10
Connecting to host 192.168.2.100, port 5201
[ 4] local 192.168.2.1 port 49762 connected to 192.168.2.100 port 5201
[ ID] Interval          Transfer      Bandwidth      Retr  Cwnd
[ 4] 0.00-1.00 sec 168 MBytes 1.41 Gbits/sec 371 328 KBytes
(omitted)
[ 4] 1.00-2.00 sec 164 MBytes 1.37 Gbits/sec 162 317 KBytes
(omitted)
[ 4] 2.00-3.00 sec 174 MBytes 1.46 Gbits/sec 56 297 KBytes
(omitted)
[ 4] 3.00-4.00 sec 167 MBytes 1.41 Gbits/sec 111 314 KBytes
(omitted)
[ 4] 4.00-5.00 sec 162 MBytes 1.36 Gbits/sec 136 291 KBytes
(omitted)

```

[4]	5.00-6.00	sec	166 MBytes	1.39 Gbits/sec	56	310 KBytes
(omitted)						
[4]	6.00-7.00	sec	171 MBytes	1.44 Gbits/sec	89	301 KBytes
(omitted)						
[4]	7.00-8.00	sec	162 MBytes	1.36 Gbits/sec	78	252 KBytes
(omitted)						
[4]	8.00-9.00	sec	163 MBytes	1.37 Gbits/sec	82	376 KBytes
(omitted)						
[4]	9.00-10.00	sec	175 MBytes	1.47 Gbits/sec	140	267 KBytes
(omitted)						
[4]	0.00-1.00	sec	164 MBytes	1.38 Gbits/sec	133	313 KBytes
[4]	1.00-2.00	sec	167 MBytes	1.40 Gbits/sec	142	287 KBytes
[4]	2.00-3.00	sec	167 MBytes	1.40 Gbits/sec	118	293 KBytes
[4]	3.00-4.00	sec	165 MBytes	1.39 Gbits/sec	84	313 KBytes
[4]	4.00-5.00	sec	166 MBytes	1.39 Gbits/sec	125	304 KBytes
[4]	5.00-6.00	sec	170 MBytes	1.42 Gbits/sec	58	356 KBytes
[4]	6.00-7.00	sec	157 MBytes	1.32 Gbits/sec	83	297 KBytes
[4]	7.00-8.00	sec	169 MBytes	1.42 Gbits/sec	97	240 KBytes
[4]	8.00-9.00	sec	164 MBytes	1.37 Gbits/sec	136	238 KBytes
[4]	9.00-10.00	sec	165 MBytes	1.38 Gbits/sec	115	242 KBytes
[4]	10.00-11.00	sec	164 MBytes	1.38 Gbits/sec	101	235 KBytes
[4]	11.00-12.00	sec	168 MBytes	1.41 Gbits/sec	72	308 KBytes
[4]	12.00-13.00	sec	174 MBytes	1.46 Gbits/sec	97	250 KBytes
[4]	13.00-14.00	sec	178 MBytes	1.49 Gbits/sec	38	337 KBytes
[4]	14.00-15.00	sec	161 MBytes	1.35 Gbits/sec	82	262 KBytes
[4]	15.00-16.00	sec	171 MBytes	1.44 Gbits/sec	53	327 KBytes
[4]	16.00-17.00	sec	169 MBytes	1.42 Gbits/sec	109	317 KBytes
[4]	17.00-18.00	sec	157 MBytes	1.32 Gbits/sec	17	366 KBytes
[4]	18.00-19.00	sec	166 MBytes	1.39 Gbits/sec	79	318 KBytes
[4]	19.00-20.00	sec	162 MBytes	1.36 Gbits/sec	154	296 KBytes
[4]	20.00-21.00	sec	172 MBytes	1.44 Gbits/sec	48	297 KBytes
[4]	21.00-22.00	sec	166 MBytes	1.39 Gbits/sec	122	283 KBytes
[4]	22.00-23.00	sec	174 MBytes	1.46 Gbits/sec	98	296 KBytes
[4]	23.00-24.00	sec	172 MBytes	1.44 Gbits/sec	75	346 KBytes
[4]	24.00-25.00	sec	167 MBytes	1.40 Gbits/sec	110	329 KBytes
[4]	25.00-26.00	sec	175 MBytes	1.46 Gbits/sec	61	321 KBytes
[4]	26.00-27.00	sec	172 MBytes	1.44 Gbits/sec	115	304 KBytes
[4]	27.00-28.00	sec	161 MBytes	1.36 Gbits/sec	83	249 KBytes
[4]	28.00-29.00	sec	171 MBytes	1.44 Gbits/sec	68	345 KBytes
[4]	29.00-30.00	sec	167 MBytes	1.41 Gbits/sec	55	356 KBytes
[4]	30.00-31.00	sec	162 MBytes	1.36 Gbits/sec	55	399 KBytes
[4]	31.00-32.00	sec	169 MBytes	1.42 Gbits/sec	103	304 KBytes
[4]	32.00-33.00	sec	161 MBytes	1.35 Gbits/sec	124	329 KBytes
[4]	33.00-34.00	sec	171 MBytes	1.43 Gbits/sec	125	322 KBytes
[4]	34.00-35.00	sec	163 MBytes	1.37 Gbits/sec	106	249 KBytes
[4]	35.00-36.00	sec	166 MBytes	1.39 Gbits/sec	115	288 KBytes
[4]	36.00-37.00	sec	161 MBytes	1.35 Gbits/sec	46	410 KBytes
[4]	37.00-38.00	sec	170 MBytes	1.43 Gbits/sec	80	283 KBytes
[4]	38.00-39.00	sec	173 MBytes	1.45 Gbits/sec	77	308 KBytes
[4]	39.00-40.00	sec	168 MBytes	1.41 Gbits/sec	79	321 KBytes
[4]	40.00-41.00	sec	166 MBytes	1.39 Gbits/sec	100	324 KBytes

```

[ 4] 41.00-42.00 sec 168 MBytes 1.41 Gbits/sec 84 331 KBytes
[ 4] 42.00-43.00 sec 155 MBytes 1.30 Gbits/sec 118 358 KBytes
[ 4] 43.00-44.00 sec 172 MBytes 1.44 Gbits/sec 95 300 KBytes
[ 4] 44.00-45.00 sec 170 MBytes 1.43 Gbits/sec 103 370 KBytes
[ 4] 45.00-46.00 sec 174 MBytes 1.46 Gbits/sec 83 320 KBytes
[ 4] 46.00-47.00 sec 163 MBytes 1.36 Gbits/sec 25 301 KBytes
[ 4] 47.00-48.00 sec 181 MBytes 1.51 Gbits/sec 111 242 KBytes
[ 4] 48.00-49.00 sec 170 MBytes 1.43 Gbits/sec 47 344 KBytes
[ 4] 49.00-50.00 sec 167 MBytes 1.40 Gbits/sec 122 342 KBytes
[ 4] 50.00-51.00 sec 165 MBytes 1.39 Gbits/sec 119 346 KBytes
[ 4] 51.00-52.00 sec 166 MBytes 1.40 Gbits/sec 129 318 KBytes
[ 4] 52.00-53.00 sec 159 MBytes 1.34 Gbits/sec 151 342 KBytes
[ 4] 53.00-54.00 sec 170 MBytes 1.42 Gbits/sec 140 331 KBytes
[ 4] 54.00-55.00 sec 172 MBytes 1.44 Gbits/sec 67 253 KBytes
^C[ 4] 55.00-55.20 sec 32.8 MBytes 1.37 Gbits/sec 0 355 KBytes
- - - - -
[ ID] Interval          Transfer      Bandwidth      Retr
[ 4] 0.00-55.20 sec 9.02 GBytes 1.40 Gbits/sec 5132
sender
[ 4] 0.00-55.20 sec 0.00 Bytes 0.00 bits/sec receiver

```

Q3.2 - On the VM, use tools/load_rules script to load filtering rules for both iptables and eBPF program. First, run it with 'ipt' argument, it creates a new (unreferenced) chain named ddos and fills it with sample IPv4 addresses stored in 10k_random_ip.txt file. Measure the performance when all the incoming traffic goes through 'ddos' chain i.e. iptables -i eth1 -I INPUT 1 -j ddos. Show that the chain is applied and is receiving traffic, also include the output of your measurements. When done, don't forget to remove the rule.

load rule:

```

root@archlinux:/vagrant/tools# ../tools/load_rules ipt
9999

```

New chain called "ddos":

```

Chain ddos (0 references)
target      prot opt source                destination            tcp
DROP        tcp  --  38.76.22.47             anywhere                tcp
DROP        tcp  --  19.221.12.198           anywhere                tcp
DROP        tcp  --  122.43.213.70           anywhere                tcp
DROP        tcp  --  81-224-221-121-no89.tbcn.telia.com anywhere
tcp
DROP        tcp  --  125.149.75.15           anywhere                tcp
DROP        tcp  --  136.31.215.92           anywhere                tcp
DROP        tcp  --  103.249.107.248         anywhere                tcp
DROP        tcp  --  167.123.94.251          anywhere                tcp
.

```

.

Now link the INPUT chain to the DDOS chain:

```
root@archlinux:/vagrant/tools# iptables -I INPUT 1 -j ddos

root@archlinux:/vagrant/tools# iptables -L
Chain INPUT (policy ACCEPT)
target     prot opt source                destination
ddos       all  --  anywhere              anywhere
```

Now all INPUT will be forwarded to DDOS chain:

```
root@archlinux:/vagrant/tools# iptables -L -v -n
Chain INPUT (policy ACCEPT 438 packets, 41898 bytes)
  pkts bytes target     prot opt in     out     source
destination
    418 40810 ddos       all  --  *      *      0.0.0.0/0
0.0.0.0/0
Chain ddos (1 references)
```

Test and check:

```
on server:
root@bristol:/home/kotaiba# sudo hping3 --rand-source 192.168.2.100 --faster

on VM:
root@archlinux:/vagrant/tools# iptables -L INPUT -v -n
Chain INPUT (policy ACCEPT 1340 packets, 78670 bytes)
  pkts bytes target     prot opt in     out     source
destination
   1340 78670 ddos       all  --  *      *      0.0.0.0/0
0.0.0.0/0
```

Now, redo the measurement:

```
root@bristol:/home/kotaiba# sudo hping3 --rand-source 192.168.2.100 --faster
&
[1] 9286
root@bristol:/home/kotaiba# HPING 192.168.2.100 (vboxnet0 192.168.2.100): NO
FLAGS are set, 40 headers + 0 data bytes

root@bristol:/home/kotaiba# iperf3 -c 192.168.2.100 -t 70 -0 10
Connecting to host 192.168.2.100, port 5201
[ 4] local 192.168.2.1 port 49772 connected to 192.168.2.100 port 5201
[ ID] Interval           Transfer     Bandwidth       Retr   Cwnd
[ 4]  0.00-1.00    sec   4.82 MBytes  40.5 Mbits/sec    78   28.3 KBytes
```



```

(omitted)
[ 4]  1.00-2.00  sec  2.34 MBytes  19.6 Mbits/sec    0   66.5 KBytes
(omitted)
[ 4]  2.00-3.00  sec  2.51 MBytes  21.1 Mbits/sec   12   26.9 KBytes
(omitted)
[ 4]  3.00-4.00  sec  4.07 MBytes  34.2 Mbits/sec    5   82.0 KBytes
(omitted)
[ 4]  4.00-5.00  sec  3.84 MBytes  32.2 Mbits/sec    7   84.8 KBytes
(omitted)
[ 4]  5.00-6.00  sec  3.82 MBytes  32.1 Mbits/sec    3   96.2 KBytes
(omitted)
[ 4]  6.00-7.00  sec  5.32 MBytes  44.6 Mbits/sec    8   94.7 KBytes
(omitted)
[ 4]  7.00-8.00  sec  5.30 MBytes  44.5 Mbits/sec    4   91.9 KBytes
(omitted)
[ 4]  8.00-9.00  sec  4.51 MBytes  37.8 Mbits/sec   47   24.0 KBytes
(omitted)
[ 4]  9.00-10.00 sec  2.72 MBytes  22.9 Mbits/sec   21   33.9 KBytes
(omitted)
[ 4]  0.00-1.00  sec  2.43 MBytes  20.4 Mbits/sec    0   69.3 KBytes
[ 4]  1.00-2.00  sec  2.14 MBytes  18.0 Mbits/sec   15   60.8 KBytes
[ 4]  2.00-3.00  sec  1.48 MBytes  12.4 Mbits/sec   11   43.8 KBytes
[ 4]  3.00-4.00  sec  2.80 MBytes  23.5 Mbits/sec    3   72.1 KBytes
[ 4]  4.00-5.00  sec  1.81 MBytes  15.2 Mbits/sec   15   41.0 KBytes
[ 4]  5.00-6.00  sec  3.83 MBytes  32.1 Mbits/sec    7   65.0 KBytes
[ 4]  6.00-7.00  sec  2.23 MBytes  18.7 Mbits/sec   13   59.4 KBytes
[ 4]  7.00-8.00  sec  6.69 MBytes  56.1 Mbits/sec    0    117 KBytes
[ 4]  8.00-9.00  sec  6.09 MBytes  51.1 Mbits/sec    4    112 KBytes
[ 4]  9.00-10.00 sec  2.30 MBytes  19.3 Mbits/sec   31   39.6 KBytes
^C[ 4] 10.00-10.62 sec   754 KBytes  9.96 Mbits/sec    3   39.6 KBytes

```

As we see a huge decrease in the bandwidth compared with previous output the speed decreased in around 1Gbit.

Now, flush and list:

```

root@archlinux:/vagrant/tools# iptables -F
root@archlinux:/vagrant/tools# iptables -L
Chain INPUT (policy ACCEPT)
target     prot opt source                destination

Chain FORWARD (policy ACCEPT)
target     prot opt source                destination

Chain OUTPUT (policy ACCEPT)
target     prot opt source                destination

Chain ddos (0 references)
target     prot opt source                destination

```

Q3.3 - Attach the eBPF object file, and fill the eBPF map using the tools/load_rules script. This time, use the 'ebpf' argument, to fill the 'ddos' bpf-map with the hex versions of the previously loaded addresses. Show that the map is applied and is receiving traffic. Repeat the previous measurement and include the output and results. Did you notice a performance difference?

```
root@archlinux:/vagrant/tools# bpf-map info /sys/fs/bpf/tc/globals/ddos
Type:      Hash
Key size:   4
Value size: 8
Max entries: 11000
Flags:      0x0
```

Execute the same performance test as in the previous section. The results are the following:

```
root@bristol:/home/kotaiba# iperf3 -c 192.168.2.100 -t 70 -0 10
Connecting to host 192.168.2.100, port 5201
[ 4] local 192.168.2.1 port 54056 connected to 192.168.2.100 port 5201
[ ID] Interval           Transfer     Bandwidth       Retr   Cwnd
[ 4]  0.00-1.00   sec    216 MBytes  1.81 Gbits/sec   213    297 KBytes
(omitted)
[ 4]  1.00-2.00   sec    195 MBytes  1.63 Gbits/sec   126    274 KBytes
(omitted)
[ 4]  2.00-3.00   sec    212 MBytes  1.78 Gbits/sec   134    339 KBytes
(omitted)
[ 4]  3.00-4.00   sec    209 MBytes  1.75 Gbits/sec   230    307 KBytes
(omitted)
[ 4]  4.00-5.00   sec    219 MBytes  1.84 Gbits/sec   333    301 KBytes
(omitted)
[ 4]  5.00-6.00   sec    205 MBytes  1.72 Gbits/sec   137    215 KBytes
(omitted)
[ 4]  6.00-7.00   sec    223 MBytes  1.87 Gbits/sec   304    250 KBytes
(omitted)
[ 4]  7.00-8.00   sec    215 MBytes  1.80 Gbits/sec   260    247 KBytes
(omitted)
[ 4]  8.00-9.00   sec    223 MBytes  1.87 Gbits/sec   266    225 KBytes
(omitted)
[ 4]  9.00-10.00  sec    214 MBytes  1.79 Gbits/sec   139    342 KBytes
(omitted)
[ 4]  0.00-1.00   sec    215 MBytes  1.80 Gbits/sec   257    352 KBytes
[ 4]  1.00-2.00   sec    221 MBytes  1.85 Gbits/sec   296    334 KBytes
[ 4]  2.00-3.00   sec    204 MBytes  1.71 Gbits/sec   151    321 KBytes
[ 4]  3.00-4.00   sec    192 MBytes  1.61 Gbits/sec   138    317 KBytes
[ 4]  4.00-5.00   sec    199 MBytes  1.67 Gbits/sec   228    305 KBytes
[ 4]  5.00-6.00   sec    221 MBytes  1.85 Gbits/sec   331    218 KBytes
[ 4]  6.00-7.00   sec    215 MBytes  1.80 Gbits/sec   167    215 KBytes
[ 4]  7.00-8.00   sec    209 MBytes  1.75 Gbits/sec   229    331 KBytes
[ 4]  8.00-9.00   sec    207 MBytes  1.74 Gbits/sec   101    291 KBytes
[ 4]  9.00-10.00  sec    193 MBytes  1.62 Gbits/sec   173    286 KBytes
```

[4]	10.00-11.00	sec	211 MBytes	1.77 Gbits/sec	333	325 KBytes
[4]	11.00-12.00	sec	216 MBytes	1.81 Gbits/sec	147	277 KBytes
[4]	12.00-13.00	sec	194 MBytes	1.63 Gbits/sec	169	229 KBytes
[4]	13.00-14.00	sec	210 MBytes	1.76 Gbits/sec	127	315 KBytes
[4]	14.00-15.00	sec	212 MBytes	1.78 Gbits/sec	231	317 KBytes
[4]	15.00-16.00	sec	212 MBytes	1.78 Gbits/sec	156	260 KBytes
[4]	16.00-17.00	sec	224 MBytes	1.88 Gbits/sec	229	335 KBytes
[4]	17.00-18.00	sec	210 MBytes	1.76 Gbits/sec	184	317 KBytes
[4]	18.00-19.00	sec	209 MBytes	1.75 Gbits/sec	171	212 KBytes
[4]	19.00-20.00	sec	216 MBytes	1.81 Gbits/sec	130	297 KBytes
[4]	20.00-21.00	sec	234 MBytes	1.96 Gbits/sec	275	240 KBytes
[4]	21.00-22.00	sec	209 MBytes	1.76 Gbits/sec	144	304 KBytes
[4]	22.00-23.00	sec	223 MBytes	1.87 Gbits/sec	177	320 KBytes
[4]	23.00-24.00	sec	212 MBytes	1.78 Gbits/sec	198	280 KBytes
[4]	24.00-25.00	sec	71.3 MBytes	640 Mbits/sec	21	205 KBytes
[4]	25.00-26.00	sec	63.9 MBytes	570 Mbits/sec	69	266 KBytes
[4]	26.00-27.00	sec	66.0 MBytes	554 Mbits/sec	42	208 KBytes
[4]	27.00-28.00	sec	65.9 MBytes	561 Mbits/sec	23	243 KBytes
[4]	28.00-29.00	sec	68.4 MBytes	565 Mbits/sec	18	201 KBytes
[4]	29.00-30.00	sec	52.1 MBytes	479 Mbits/sec	38	235 KBytes
[4]	30.00-31.00	sec	61.6 MBytes	550 Mbits/sec	60	266 KBytes
[4]	31.00-32.00	sec	57.7 MBytes	484 Mbits/sec	18	204 KBytes
[4]	32.00-33.00	sec	231 MBytes	1.75 Gbits/sec	231	290 KBytes
[4]	33.00-34.00	sec	201 MBytes	1.75 Gbits/sec	201	290 KBytes
[4]	34.00-35.00	sec	208 MBytes	1.75 Gbits/sec	219	290 KBytes

As we see above I used the hping again and we can notice the effect of eBPF on bandwidth is much smaller than the iptables. In addition to eBPF shows speeds very similar to the raw speeds, while iptables has lower speeds.

Task 4: Understanding eBPF architecture (2 points)

The answers for this task is based on bpf man page as stated in the source.

Q4.1 - What is the difference between (c)BPF and eBPF?

Extended BPF (or eBPF) is similar to the original ("classic") BPF (cBPF) used to filter network packets. eBPF extends cBPF in multiple ways, including the ability to call a fixed set of in-kernel helper functions (via the BPF_CALL opcode extension provided by eBPF) and access shared data structures such as eBPF maps.

Q4.2 - What is the purpose of eBPF maps?

eBPF maps are a generic data structure for storage of different data types. Data types are generally treated as binary blobs, so a user just specifies the size of the key and the size of the value at map-creation time. In other words, a key/value for a given map can have an arbitrary structure.

Q4.3 - How does an eBPF program gets passed to the kernel? Which user-level tools are used for this?

First, the kernel statically analyzes the programs before loading them, in order to ensure that they cannot harm the running system. Then, the BPF programs are loaded into the kernel by the `bpf()` system call. The user `tc-bpf` tools used to remains alive so that the `tc`-system knows which module specifically to load.

Q4.4 - What kind of operations can be performed on a network packet inside eBPF code?

A possible kind of operations that can be performed are network packets can be forwarded, dropped and changed. We can adjust and modify every byte in a network packet in a way that the network packet will change to something totally different based on changes that are pre-programmed conditions. In general, BPF itself decides whether to drop or forward packets to which destination.

Source:

<http://man7.org/linux/man-pages/man2/bpf.2.html>